

432 Homework 1 Answer Sketch

Thomas E. Love

Due 2018-01-26. Version: 2018-01-29

Contents

0.1	Setup and Data Ingest	1
1	Question 1 (40 points)	2
1.1	Creating the BMI and BMI Category variables	2
1.2	Checking for Missing Values	3
1.3	First Attempt at Table 1	5
1.4	Making Decisions about Summary Approaches	5
1.5	Final Table 1 and Notes	10
2	Question 2. (30 points)	11
2.1	We'll start with a summary table and graph	11
2.2	A Two-Way ANOVA model with Interaction	13
2.3	A Two-Way ANOVA model without interaction (Main Effects only)	14
3	Question 3 (25 points)	14
3.1	ANCOVA: Adding age to the "no interaction" model	14
3.2	Comparison to the "no interaction" model via ANOVA	14
3.3	Comparison to the "no interaction" model via fit quality summaries	15
3.4	Comparison to the "no interaction" model in terms of conclusions	15
3.5	Regression diagnostics for the ANCOVA model	16
4	Question 4. (5 points)	16

0.1 Setup and Data Ingest

Note: It's good practice not to load any more than you must, although you should install the complete list of packages on our web site.

```
library(skimr)
library(broom)
library(tableone)
library(tidyverse)
```

```
-- Attaching packages -----
v ggplot2 2.2.1    v purrr   0.2.4
v tibble  1.4.2    v dplyr   0.7.4
v tidyr   0.7.2    v stringr 1.2.0
v readr   1.1.1    v forcats 0.2.0

-- Conflicts -----
x dplyr::contains() masks skimr::contains()
x dplyr::ends_with() masks skimr::ends_with()
x dplyr::everything() masks skimr::everything()
x dplyr::filter()    masks stats::filter()
x dplyr::lag()        masks stats::lag()
```

```
x dplyr::matches()      masks skimr::matches()
x dplyr::num_range()    masks skimr::num_range()
x dplyr::one_of()       masks skimr::one_of()
x dplyr::starts_with() masks skimr::starts_with()

skim_with(numeric = list(hist = NULL), integer = list(hist = NULL))
```

Note: I loaded the data for this assignment into a subfolder of my R Project directory for Homework 1 called `data`. Hence, I use the following command to load in the `hbp330.csv` data.

```
hbp330 <- read.csv("data/hbp330.csv") %>% tbl_df
```

1 Question 1 (40 points)

Note: I don't expect or need you to include the question in your response to homework assignments. I include them here to provide some context for the sketch.

Build a Table 1 to compare the subjects in practice A to the subjects in practice B on the following nine variables: age, race, Hispanic ethnicity, sex, primary insurance, body mass index, BMI category, and systolic and diastolic blood pressure. Make the Table as well as you can within R, and display the result as part of your HTML file. Include a description of the results of your Table 1 that does not exceed 100 words, using complete English sentences.

1.1 Creating the BMI and BMI Category variables

With the metric system, the formula for BMI is weight (in kilograms) divided by the square of height (in meters.)

```
hbp330 <- hbp330 %>%

  # first we'll create the bmi values

  mutate( bmi = weight / (height*height) ) %>%

  # next we'll create the bmi categories with case_when
  # note that values of NA in bmi return NA in bmi_cat, too

  mutate( bmi_cat = case_when(
    bmi < 18.5 ~ "Underweight",
    bmi < 25 ~ "Normal",
    bmi < 30 ~ "Overweight",
    bmi >= 30 ~ "Obese")) %>%

  # finally we arrange the bmi_cat by the median(bmi) levels
  # we could instead have arranged by hand using fct_relevel

  mutate( bmi_cat = fct_reorder( bmi_cat, bmi, median ))
```

As a sanity check, let's ensure that the subjects in each `bmi_cat` have bmi values as we expect...

```
hbp330 %>% group_by(bmi_cat) %>% skim(bmi)
```

```
Skim summary statistics
n obs: 330
```

```
n variables: 21
group variables: bmi_cat
```

Variable type: numeric

	bmi_cat	variable	missing	complete	n	mean	sd	p0	p25	median
Underweight	bmi	0	2	2	17.41	0.97	16.73	17.07	17.41	
Normal	bmi	0	25	25	22.96	1.67	19.68	21.41	23.68	
Overweight	bmi	0	63	63	27.66	1.43	25.05	26.6	27.32	
Obese	bmi	0	240	240	38.1	6.81	30.01	33.06	36.67	
p75	p100									
17.76	18.1									
24.32	24.99									
28.88	30									
40.97	64.04									

OK. The minima and maxima match what we're looking for.

Note: In practical work, we would likely collapse together the Underweight and Normal categories, since there are so few patients in the Underweight category.

```
hbp330 <- hbp330 %>%
  mutate(bmi_cat3 = fct_collapse(bmi_cat,
                                "Not_Overweight" = c("Underweight", "Normal")))
```

And again, a little sanity check to ensure we've not made a mistake.

```
hbp330 %>% count(bmi_cat, bmi_cat3)
```

```
# A tibble: 4 x 3
  bmi_cat  bmi_cat3      n
  <fct>    <fct>    <int>
1 Underweight Not_Overweight      2
2 Normal      Not_Overweight     25
3 Overweight  Overweight      63
4 Obese       Obese          240
```

1.2 Checking for Missing Values

```
hbp330 %>%
  group_by(practice) %>%
  skim()
```

```
Skim summary statistics
n obs: 330
n variables: 22
group variables: practice
```

Variable type: factor

practice	variable	missing	complete	n	n_unique
A	bmi_cat	0	180	180	4
A	bmi_cat3	0	180	180	3
A	depdiag	0	180	180	2
A	eth_hisp	4	176	180	2
A	insurance	0	180	180	4
A	provider	0	180	180	7

A	race	1	179	180	3
A	sex	0	180	180	2
A	subject	0	180	180	180
A	tobacco	0	180	180	3
B	bmi_cat	0	150	150	4
B	bmi_cat3	0	150	150	3
B	depdiag	0	150	150	2
B	eth_hisp	1	149	150	2
B	insurance	0	150	150	4
B	provider	0	150	150	7
B	race	1	149	150	4
B	sex	0	150	150	2
B	subject	0	150	150	150
B	tobacco	0	150	150	3

top_counts ordered

Obe: 136, Ove: 32, Nor: 11, Und: 1 FALSE

Obe: 136, Ove: 32, Not: 12, NA: 0 FALSE

No: 102, Yes: 78, NA: 0 FALSE

No: 174, NA: 4, Yes: 2 FALSE

Med: 76, Med: 66, Com: 35, Uni: 3 FALSE

A10: 44, A10: 37, A10: 32, A10: 31 FALSE

Bla: 166, Whi: 9, Mul: 4, NA: 1 FALSE

F: 119, M: 61, NA: 0 FALSE

A00: 1, A00: 1, A00: 1, A00: 1 FALSE

for: 75, nev: 70, cur: 35, NA: 0 FALSE

Obe: 104, Ove: 31, Nor: 14, Und: 1 FALSE

Obe: 104, Ove: 31, Not: 15, NA: 0 FALSE

No: 112, Yes: 38, NA: 0 FALSE

No: 87, Yes: 62, NA: 1 FALSE

Med: 68, Med: 54, Com: 18, Uni: 10 FALSE

B10: 29, B10: 25, B10: 25, B10: 19 FALSE

Whi: 122, Bla: 14, Asi: 10, Mul: 3 FALSE

F: 84, M: 66, NA: 0 FALSE

B00: 1, B00: 1, B00: 1, B00: 1 FALSE

nev: 70, for: 42, cur: 38, NA: 0 FALSE

Variable type: integer

practice	variable	missing	complete	n	mean	sd	p0	p25
A	age	0	180	180	56.34	11.17	24	49
A	bpmed	0	180	180	0.64	0.48	0	0
A	dbp	0	180	180	74.49	11.4	41	67
A	hsgrad	0	180	180	82.73	12.48	-2	78
A	income	0	180	180	34780	19412.55	100	24675
A	ldl	27	153	180	108.18	37.16	39	82
A	sbp	0	180	180	130.82	15.38	84	120
A	statin	0	180	180	0.7	0.46	0	0
B	age	0	150	150	54.17	11.89	23	46
B	bpmed	0	150	150	0.67	0.47	0	0
B	dbp	0	150	150	75.05	8.58	54	69
B	hsgrad	0	150	150	80.03	7.72	68	74
B	income	0	150	150	35799.33	10768.65	15600	25925
B	ldl	51	99	150	96.06	33.85	35	70
B	sbp	0	150	150	125.44	19	85	113
B	statin	0	150	150	0.71	0.45	0	0

median	p75	p100
58	65	77
1	1	1
74	83	101
84	91	100
29150	42825	147400
102	132	245
131	139	194
1	1	1
55.5	64	75
1	1	1
74.5	80	106
78	86.75	97
34550	42275	71400
93	123.5	174
123.5	134.75	191
1	1	1

Variable type: numeric

practice	variable	missing	complete	n	mean	sd	p0	p25	median
A	bmi	0	180	180	35.2	8.2	16.73	30.05	34.06
A	height	0	180	180	1.68	0.12	1.38	1.6	1.67
A	weight	0	180	180	98.51	22.76	53	82.52	94.5
B	bmi	0	150	150	34.39	7.83	18.1	28.74	33.48
B	height	0	150	150	1.64	0.12	1.4	1.55	1.63
B	weight	0	150	150	92.39	23.42	50.4	75.6	90.45
	p75	p100							
	39.3	64.04							
	1.76	1.95							
	112.53	162.4							
	38.98	63.55							
	1.71	1.93							
	106.97	181.5							

We're missing the following values:

- `ldl` on 78 subjects, but that's not in our Table 1, so we won't worry about that now, and
- `race` on 2 subjects (1 in practice A and 1 in practice B), and
- `eth_hisp` on 5 subjects (4 in practice A and 1 in practice B).

1.3 First Attempt at Table 1

```
hw1_table1 <- CreateTableOne(data = hbp330,
  vars = c("age", "race", "eth_hisp", "sex",
    "insurance", "bmi", "bmi_cat",
    "sbp", "dbp"),
  strata = c("practice"))
```

1.4 Making Decisions about Summary Approaches

Note that if we look at a `summary` of this Table 1, we see that:

- the difference in p values between the “normal” and “non-normal” versions of each of the continuous variables is small enough that if we’re going to compare the p values to, say, 0.05, it won’t matter which we choose, and
- the practical impact of the choice of p values between the “approximate” and “exact” versions of each of the categorical variables is also small.

```
summary(hw1_table1)
```

```
### Summary of continuous variables ###
```

```
practice: A
```

	n	miss	p.miss	mean	sd	median	p25	p75	min	max	skew	kurt
age	180	0	0	56	11	58	49	65	24	77	-0.57	-0.03
bmi	180	0	0	35	8	34	30	39	17	64	0.93	1.35
sbp	180	0	0	131	15	131	120	139	84	194	0.57	1.89
dbp	180	0	0	74	11	74	67	83	41	101	-0.07	-0.20

```
practice: B
```

	n	miss	p.miss	mean	sd	median	p25	p75	min	max	skew	kurt
age	150	0	0	54	12	56	46	64	23	75	-0.4	-0.4
bmi	150	0	0	34	8	33	29	39	18	64	0.8	1.3
sbp	150	0	0	125	19	124	113	135	85	191	0.7	0.6
dbp	150	0	0	75	9	74	69	80	54	106	0.2	0.7

```
p-values
```

	pNormal	pNonNormal
age	0.088489264	0.1122326040
bmi	0.365264000	0.4971154515
sbp	0.004781282	0.0003634655
dbp	0.617464601	0.7253555322

```
Standardize mean differences
```

```
1 vs 2
age 0.18834212
bmi 0.10044371
sbp 0.31105472
dbp 0.05596318
```

```
### Summary of categorical variables ###
```

```
practice: A
```

var	n	miss	p.miss	level	freq	percent	cum.percent
race	180	1	0.6	Asian/PI	0	0.0	0.0
				Black/AA	166	92.7	92.7
				Multi-Racial	4	2.2	95.0
				White	9	5.0	100.0
eth_hisp	180	4	2.2	No	174	98.9	98.9
				Yes	2	1.1	100.0
sex	180	0	0.0	F	119	66.1	66.1
				M	61	33.9	100.0

insurance	180	0	0.0	Commercial	35	19.4	19.4
				Medicaid	66	36.7	56.1
				Medicare	76	42.2	98.3
				Uninsured	3	1.7	100.0
bmi_cat	180	0	0.0	Underweight	1	0.6	0.6
				Normal	11	6.1	6.7
				Overweight	32	17.8	24.4
				Obese	136	75.6	100.0

practice: B

var	n	miss	p.miss	level	freq	percent	cum.percent
race	150	1	0.7	Asian/PI	10	6.7	6.7
				Black/AA	14	9.4	16.1
				Multi-Racial	3	2.0	18.1
				White	122	81.9	100.0
eth_hisp	150	1	0.7	No	87	58.4	58.4
				Yes	62	41.6	100.0
sex	150	0	0.0	F	84	56.0	56.0
				M	66	44.0	100.0
insurance	150	0	0.0	Commercial	18	12.0	12.0
				Medicaid	68	45.3	57.3
				Medicare	54	36.0	93.3
				Uninsured	10	6.7	100.0
bmi_cat	150	0	0.0	Underweight	1	0.7	0.7
				Normal	14	9.3	10.0
				Overweight	31	20.7	30.7
				Obese	104	69.3	100.0

p-values

	pApprox	pExact
race	1.041050e-50	7.529447e-61
eth_hisp	2.207408e-19	1.531820e-22
sex	7.738092e-02	6.927716e-02
insurance	1.593679e-02	1.580010e-02
bmi_cat	5.868027e-01	5.894708e-01

Standardize mean differences

	1 vs 2
race	3.1557762
eth_hisp	1.1352992
sex	0.2084780
insurance	0.3569322
bmi_cat	0.1531995

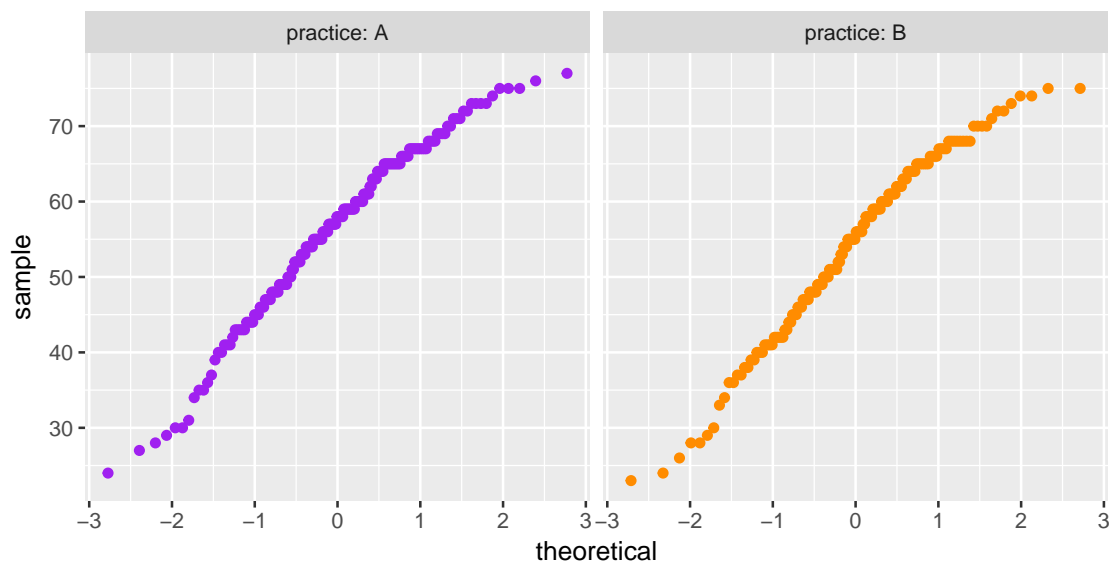
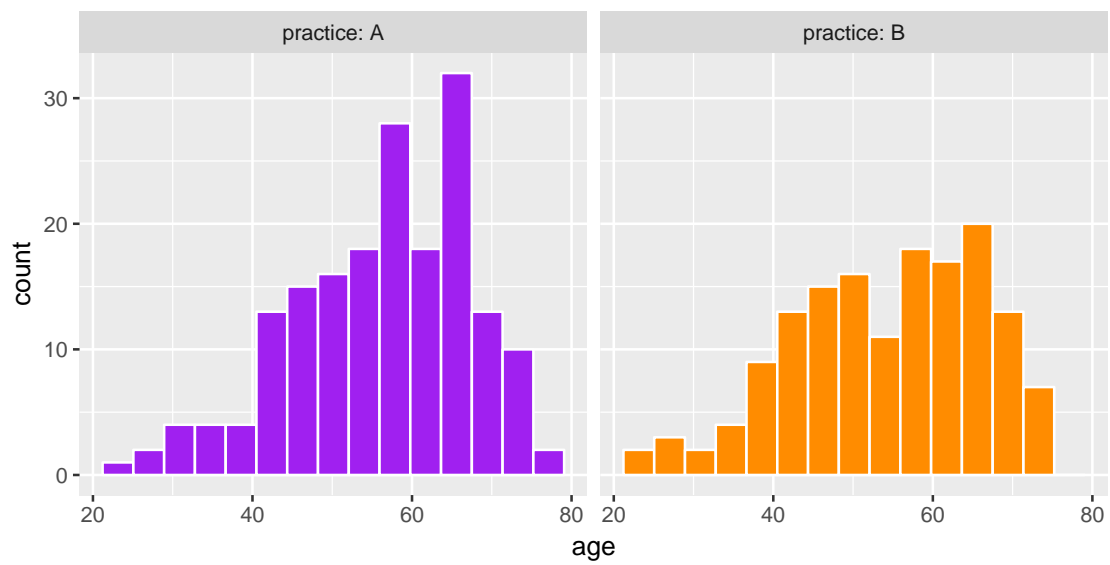
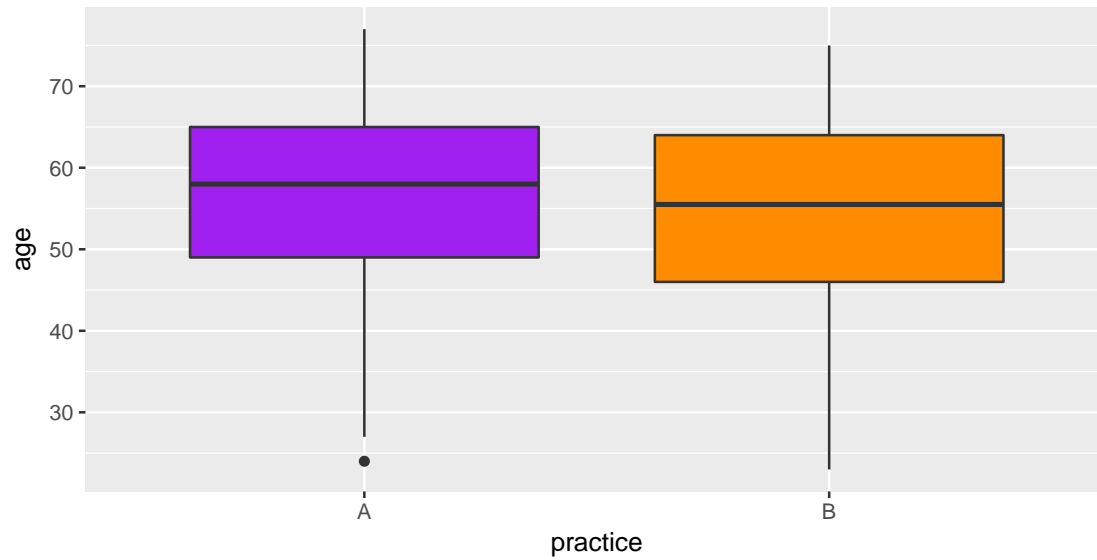
1.4.1 A closer look at Age

As an example, let's plot the `age` data within each practice to evaluate Normality. Just for fun, I'll run

- a boxplot
- a histogram, and
- a normal Q-Q plot

of the ages within each practice.

```
p1 <- ggplot(hbp330, aes(x = practice, y = age)) +  
  geom_boxplot(fill = c("purple", "darkorange"))  
  
p2 <- ggplot(hbp330, aes(x = age, fill = practice)) +  
  geom_histogram(bins = 15, col = "white") +  
  scale_fill_manual(values = c("purple", "darkorange")) +  
  guides(fill = FALSE) +  
  facet_wrap( ~ practice, labeller = label_both)  
  
p3 <- ggplot(hbp330, aes(sample = age, color = practice)) +  
  geom_qq() +  
  scale_color_manual(values = c("purple", "darkorange")) +  
  guides(color = FALSE) +  
  facet_wrap( ~ practice, labeller = label_both)  
  
gridExtra::grid.arrange(p1, p2, p3, ncol=1)
```

The data aren't *perfectly* Normally distributed, of course, but I don't see a big problem with summarizing via means and standard deviations. We could do something similar with each of the other quantitative variables.

1.5 Final Table 1 and Notes

```
print(hw1_table1)
```

	Stratified by practice			
	A	B	p	test
n	180	150		
age (mean (sd))	56.34 (11.17)	54.17 (11.89)	0.088	
race (%)			<0.001	
Asian/PI	0 (0.0)	10 (6.7)		
Black/AA	166 (92.7)	14 (9.4)		
Multi-Racial	4 (2.2)	3 (2.0)		
White	9 (5.0)	122 (81.9)		
eth_hisp = Yes (%)	2 (1.1)	62 (41.6)	<0.001	
sex = M (%)	61 (33.9)	66 (44.0)	0.077	
insurance (%)			0.016	
Commercial	35 (19.4)	18 (12.0)		
Medicaid	66 (36.7)	68 (45.3)		
Medicare	76 (42.2)	54 (36.0)		
Uninsured	3 (1.7)	10 (6.7)		
bmi (mean (sd))	35.20 (8.20)	34.39 (7.83)	0.365	
bmi_cat (%)			0.587	
Underweight	1 (0.6)	1 (0.7)		
Normal	11 (6.1)	14 (9.3)		
Overweight	32 (17.8)	31 (20.7)		
Obese	136 (75.6)	104 (69.3)		
sbp (mean (sd))	130.82 (15.38)	125.44 (19.00)	0.005	
dbp (mean (sd))	74.49 (11.40)	75.05 (8.58)	0.617	

Notes for Table 1:

1. There are 4 subjects missing Hispanic ethnicity status in practice A, and 1 in practice B.
2. There is 1 subject in each practice missing Race.
3. Results are shown in terms of means and standard deviations for quantitative variables, and t tests are used for comparisons, because a Normal approximation was a reasonable choice for each such variable.
4. For categorical variables, we display counts and percentages, and use Pearson chi-square tests of significance.

1.5.1 Describing the Table in a few sentences

The key conclusions of this Table 1 are that the two practices are not especially comparable on demographic features, or systolic blood pressure, but show generally similar diastolic blood pressure and body mass index distributions. Subjects seen in Practice A are more frequently of Black race, less often of Hispanic ethnicity, more likely to have Commercial or Medicare insurance (but less likely to have Medicaid or be uninsured) and display larger systolic blood pressure values than Practice B subjects. Practice A patients are also a bit older on average, and more likely to be female.

2 Question 2. (30 points)

Does which practice a person attends seem to have a meaningful impact on their systolic blood pressure, adjusting for whether or not they are on a blood pressure medication? Decide whether your model should include an interaction term sensibly, and then fit your choice of model and interpret and display the findings carefully. Be sure to provide a written explanation of your findings, in complete sentences. Responses without graphs are not complete.

2.1 We'll start with a summary table and graph

We want to understand something about the impact of practice on SBP, adjusting for BP medication status.

Let's start with some numbers. We'll run a count of the number of patients within the practice and medication groups, along with some basic summaries of the `sbp` data within each group.

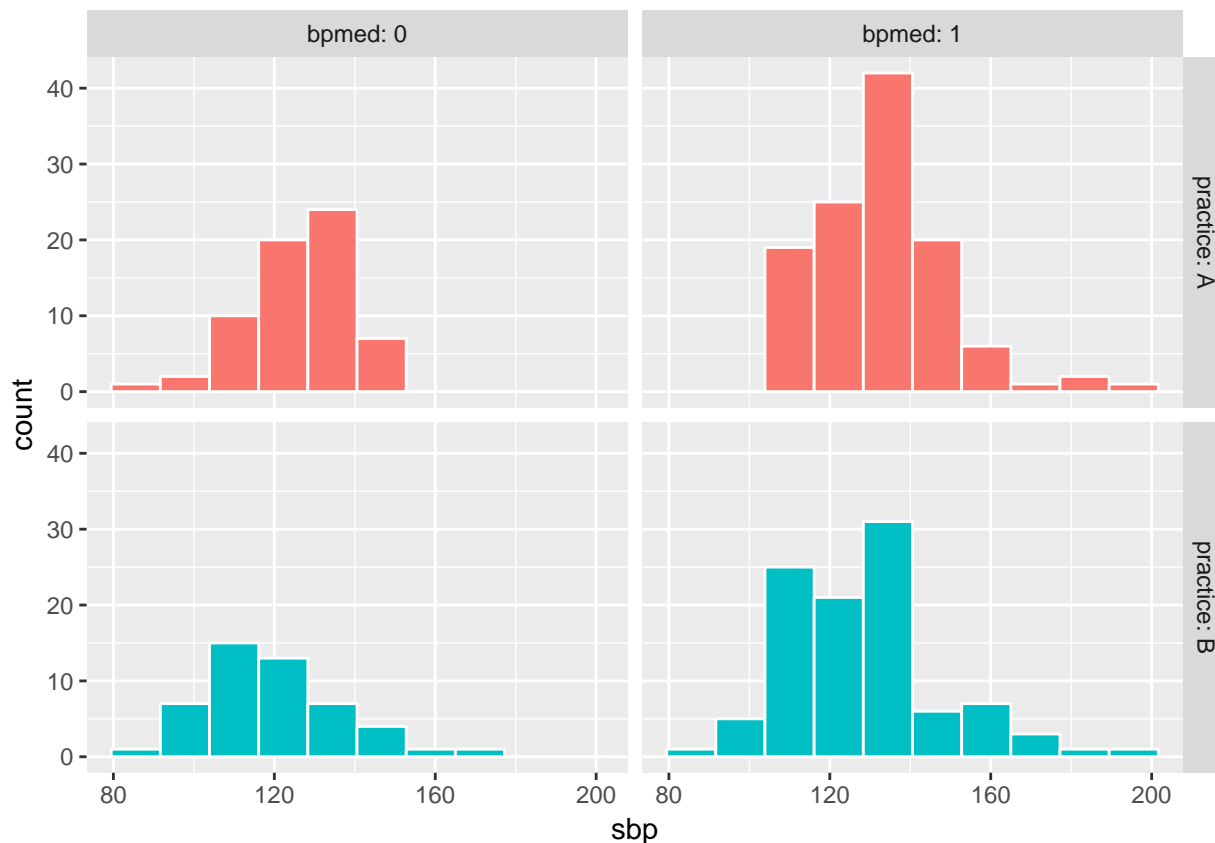
```
hbp330_summaries <- hbp330 %>%  
  group_by(practice, bpmed) %>%  
  summarize(n = n(), median = median(sbp), mean = mean(sbp), stdev = sd(sbp))
```

```
hbp330_summaries
```

```
# A tibble: 4 x 6  
# Groups:   practice [?]  
  practice bpmed      n median  mean stdev  
  <fct>    <int> <int>  <dbl> <dbl> <dbl>  
1 A            0    64   128   127  13.3  
2 A            1   116   133   133  16.0  
3 B            0    49   117   119  17.4  
4 B            1   101   127   128  19.1
```

It looks like we have a plausible sample size to look at a picture in each case (there are no very small cells here) and that the medians and means are quite close in each group. Let's try some graphs, first of the raw data distributions...

```
ggplot(hbp330, aes(x = sbp, fill = factor(practice))) +  
  geom_histogram(bins = 10, col = "white") +  
  guides(fill = FALSE) +  
  facet_grid(practice ~ bpmed, labeller = label_both)
```

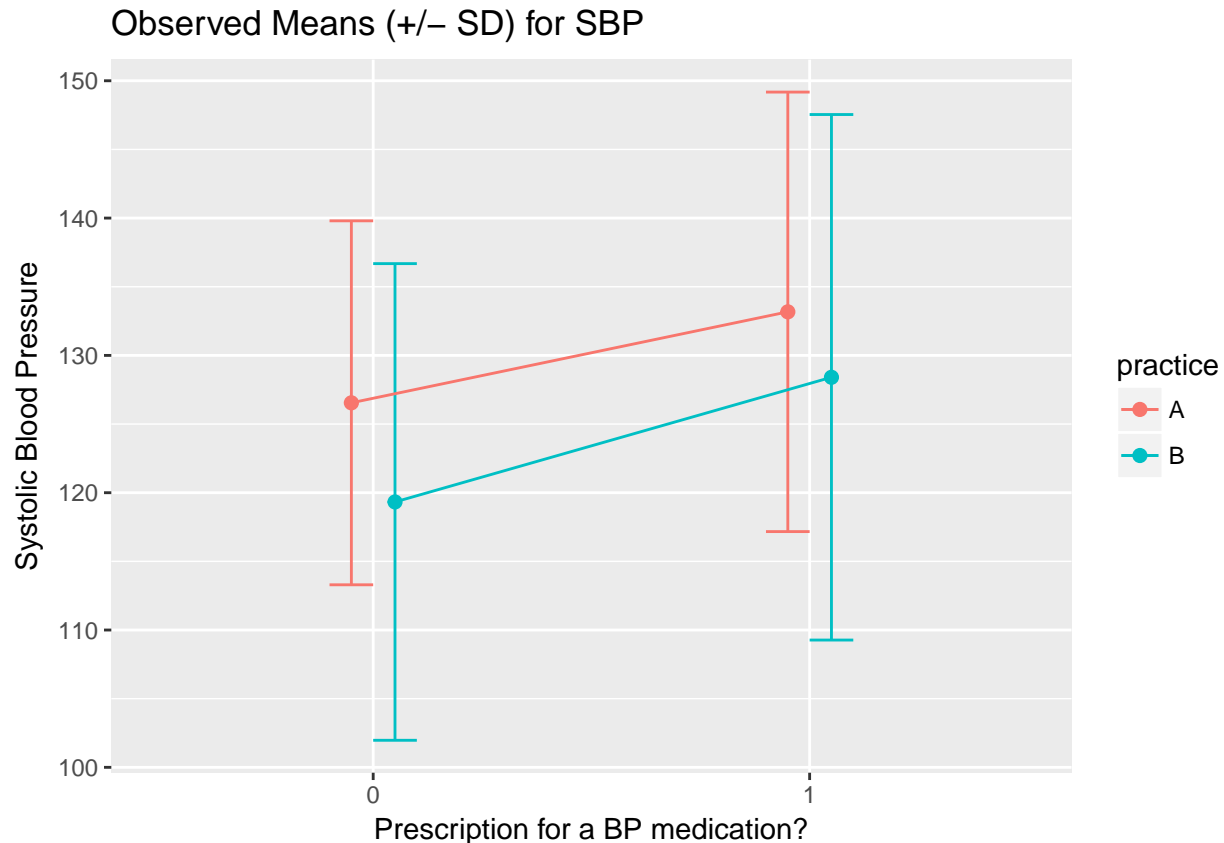


and I don't see any major problems with assuming that a mean and standard deviation might be reasonable choices to summarize the data from these four groups.

So, let's try graphing the means and standard deviations...

```
pd <- position_dodge(0.2)

ggplot(hbp330_summaries, aes(x = factor(bpmed), y = mean, col = practice)) +
  geom_errorbar(aes(ymin = mean - stdev,
                    ymax = mean + stdev),
                width = 0.2, position = pd) +
  geom_point(size = 2, position = pd) +
  geom_line(aes(group = practice), position = pd) +
  labs(y = "Systolic Blood Pressure",
       x = "Prescription for a BP medication?",
       title = "Observed Means (+/- SD) for SBP")
```



I don't see much to suggest a meaningful interaction here. The lines joining the points are essentially parallel. It looks like the group with the lowest (healthiest) mean SBP are the subjects in practice B without a medication.

2.2 A Two-Way ANOVA model with Interaction

We'll run the two models (with and without interaction) but I expect the interaction to play only a small role.

```
hw1_q2_with_int <- lm(sbp ~ practice*bpmed, data = hbp330)
anova(hw1_q2_with_int)
```

Analysis of Variance Table

Response: sbp

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
practice	1	2365	2365.2	8.4176	0.003969 **
bpmed	1	4420	4420.0	15.7303	8.984e-05 ***
practice:bpmed	1	110	110.4	0.3928	0.531264
Residuals	326	91602	281.0		

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

It doesn't look like the interaction is either significant, or accounts for a large fraction of the variation in the SBP values. So we'll run and interpret the ANOVA model without interaction.

2.3 A Two-Way ANOVA model without interaction (Main Effects only)

```
hw1_q2_no_int <- lm(sbp ~ practice + bpmed, data = hbp330)
summary(hw1_q2_no_int)
```

Call:

```
lm(formula = sbp ~ practice + bpmed, data = hbp330)
```

Residuals:

Min	1Q	Median	3Q	Max
-41.844	-11.961	-0.702	9.369	63.039

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	125.844	1.768	71.164	< 2e-16 ***
practiceB	-5.600	1.852	-3.023	0.0027 **
bpmed	7.716	1.944	3.970	8.85e-05 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 16.75 on 327 degrees of freedom

Multiple R-squared: 0.06889, Adjusted R-squared: 0.06319

F-statistic: 12.1 on 2 and 327 DF, p-value: 8.548e-06

Since each of the two factors is binary, we can simply read off that both `practice` and `bpmed` appear to have a significant impact on SBP, with practice B having lower SBP levels, on average, and subjects without BP medications having lower SBP levels, on average.

3 Question 3 (25 points)

Does the addition of the subject's age add meaningful predictive value to the model you developed in question 2? Does it change the nature of the conclusions you can draw from the model? How do you know? Be sure to provide a written explanation of your findings, in complete sentences, to accompany any output you choose to include. Responses without graphs are not complete.

3.1 ANCOVA: Adding age to the “no interaction” model

We'll build the new (ANCOVA) model including `age` and compare it to the no-interaction two-way ANOVA model.

```
hw1_q3 <- lm(sbp ~ practice + bpmed + age, data = hbp330)
```

3.2 Comparison to the “no interaction” model via ANOVA

```
anova(hw1_q3, hw1_q2_no_int)
```

Analysis of Variance Table

Model 1: sbp ~ practice + bpmed + age

```
Model 2: sbp ~ practice + bpmed
      Res.Df  RSS Df Sum of Sq    F Pr(>F)
1      326 90178
2      327 91712 -1    -1534.3 5.5467 0.01911 *
```

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

It does appear that age adds significant predictive value to the no-interaction model.

3.3 Comparison to the “no interaction” model via fit quality summaries

In terms of summarizing the two models,

```
glance(hw1_q3)
```

```
      r.squared adj.r.squared    sigma statistic    p.value df    logLik
1 0.08446494    0.07603977 16.63185  10.02531 2.452493e-06  4 -1393.973
      AIC      BIC deviance df.residual
1 2797.946 2816.941  90177.6         326
```

```
glance(hw1_q2_no_int)
```

```
      r.squared adj.r.squared    sigma statistic    p.value df    logLik
1 0.06888764    0.06319276 16.74708  12.09642 8.547555e-06  3 -1396.757
      AIC      BIC deviance df.residual
1 2801.513 2816.71  91711.92         327
```

The model with age included performs a bit better in terms of adjusted (and raw) R^2 and AIC and performs comparably in terms of BIC.

3.4 Comparison to the “no interaction” model in terms of conclusions

```
summary(hw1_q3)
```

Call:

```
lm(formula = sbp ~ practice + bpmed + age, data = hbp330)
```

Residuals:

```
      Min       1Q   Median       3Q      Max
-42.670 -11.420  -1.185   9.917  59.559
```

Coefficients:

```
              Estimate Std. Error t value Pr(>|t|)
(Intercept) 115.65014    4.67108  24.759 < 2e-16 ***
practiceB    -5.14834    1.84952  -2.784 0.005689 **
bpmed         6.63315    1.98433   3.343 0.000926 ***
age           0.19333    0.08209   2.355 0.019108 *
---
```

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 16.63 on 326 degrees of freedom

Multiple R-squared: 0.08446, Adjusted R-squared: 0.07604

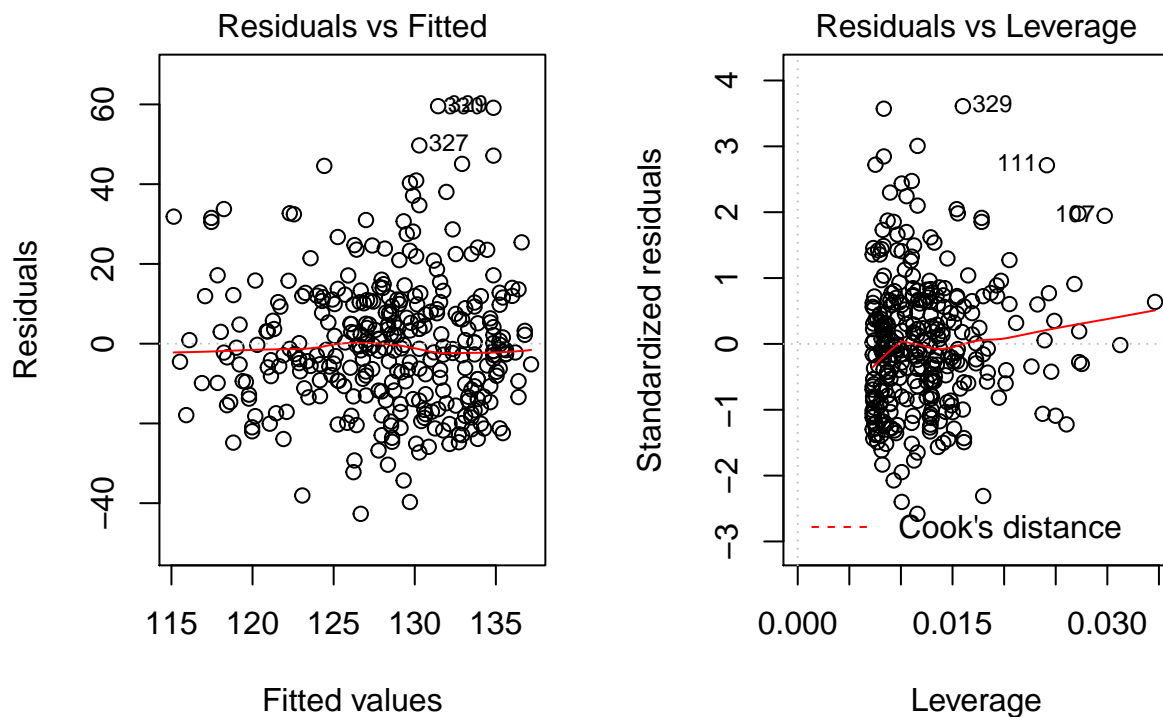
F-statistic: 10.03 on 3 and 326 DF, p-value: 2.452e-06

In terms of conclusions from the model, the adjustment for **age** does not change the fundamental conclusions from the no-interaction model. Specifically, we continue to see a significant effect of both **practice** (with B showing lower SBPs) and **bpmed** (with those not prescribed such a medication having smaller SBPs.)

3.5 Regression diagnostics for the ANCOVA model

I don't see any substantial problems with regression residuals in either a plot of residuals vs. fitted values, and there are no particularly influential points.

```
par(mfrow = c(1,2))
plot(hw1_q3, which = c(1, 5))
```



```
par(mfrow = c(1,1))
```

4 Question 4. (5 points)

Provide your GitHub name.

No sketch here. We want to know your GitHub name because we're hoping that before the semester is over, we can figure out a way to facilitate your posting the best and most sharable parts of your project portfolios to GitHub rather than just to Canvas. But for now, we just need the list of GitHub names.