

432 Class 14 Slides

github.com/THOMASELOVE/432-2018

2018-03-01

Setup

```
library(skimr); library(MASS)
library(robustbase); library(quantreg)
library(lmtest); library(sandwich)
library(boot); library(broom)
library(rms)
library(tidyverse)

decim <- function(x, k) format(round(x, k), nsmall=k)
```

Today's Materials

- Crime in the United States
- Sandwich Estimation of Standard Errors
- Bootstrapping Regression Coefficients

Next Time

- 1 Robust Linear Regression Methods with Huber weights
- 2 Robust Linear Regression with bisquare weights (biweights)
- 3 Bounded Influence Regression & Least Trimmed Squares
- 4 Penalized Least Squares using `ols` in the `rms` package
- 5 Quantile Regression on the Median

Some Motivating Graphs

A Simple Regression Model

Suppose we were looking at a simple regression on a new batch of data.


```
set.seed(20170421)
newd <- data_frame(x = 1:18, y = rnorm(x, mean = x))
newd$y[2] <- 24

head(newd)
```

```
# A tibble: 6 x 2
```

	x	y
	<int>	<dbl>
1	1	1.66
2	2	24.0
3	3	4.37
4	4	4.78
5	5	4.61
6	6	7.33

Scatterplot of newd




Code for Last Slide and Next Slide

```
(p <- ggplot(newd, aes(x = x, y = y)) +  
  geom_point(size = 3) +  
  theme_bw()  
)
```

Add OLS line

```
p + geom_smooth(method = "lm", col = "black")
```

OLS regression line




That outlier seems like a problem.

Suppose we compare the ordinary least squares regression line we saw above to a new line, fit without including the outlier at the top left of the plot.

Code for next plot

```
p + geom_smooth(method = "lm", se = FALSE, col = "black") +  
  geom_smooth(method = "lm", data = filter(newd, y < 20),  
              col = "purple", linetype = "dashed")
```

New Plot showing the outlier effect




Add robust regression lines

Now, let's add a line from a robust regression [the robust (Huber weights) line] with the “rlm” method, and a quantile regression line using the “rq” method.

- Linear model (OLS) will be in black
- Linear model (OLS) without the outlier in dashed purple
- Robust Linear Model (via `rlm`) in blue
- Quantile Regression Model (via `rq`) in red

Comparison of models



Comparison of models (code)

```
p + geom_smooth(method = "lm", col = "black", se = FALSE) +  
  geom_smooth(method = "lm", data = filter(newd, y < 20),  
              col = "purple", se = FALSE,  
              linetype = "dashed") +  
  geom_smooth(method = "rlm", col = "blue", se = FALSE) +  
  geom_smooth(method = "rq", col = "red", se = FALSE)
```

Sources and Resources

Key sources for this document include:

- <http://stats.idre.ucla.edu/r/dae/robust-regression/>
- <http://www.alastairsanderson.com/R/tutorials/robust-regression-in-R/>
- John Fox's appendix on Applied Robust Regression
- <https://cran.r-project.org/web/packages/robust/robust.pdf>
- <https://cran.r-project.org/web/packages/rms/rms.pdf>
- <http://www.statmethods.net/advstats/bootstrapping.html>

The crimestat Data

Data Source

The crimestat data gathered here refer to 2016, mainly, and were obtained from:

- <http://www.worldatlas.com/articles/the-most-dangerous-states-in-the-u-s.html>
- <https://www.statista.com/statistics/242302/percentage-of-single-mother-households-in-the-us-by-state/>
- and a few different Wikipedia sites,

but the use of these data in this context is due to an older data set that appears in *Statistical Methods for Social Sciences*, Third Edition by Alan Agresti and Barbara Finlay (Prentice Hall, 1997), and which is the primary example at <http://stats.idre.ucla.edu/r/dae/robust-regression/>

The crimestat data set

For each of 51 states (including the District of Columbia), we have the state's ID number, postal abbreviation and full name, as well as:

- **crime** - the violent crime rate per 100,000 people
- **poverty** - the official poverty rate (% of people living in poverty in the state/district) in 2014
- **single** - the percentage of households in the state/district led by a female householder with no spouse present and with her own children under 18 years living in the household in 2016
- **trump** - whether Donald Trump won the popular vote in the 2016 presidential election in that state/district (which we'll ignore for today)

The crimestat data set

```
crimestat <- read.csv("crimestat.csv") %>% tbl_df(  
  crimestat
```

```
# A tibble: 51 x 7
```

	sid	state	crime	poverty	single	trump	state.full
	<int>	<fct>	<dbl>	<dbl>	<dbl>	<int>	<fct>
1	1	AL	427	19.2	9.02	1	Alabama
2	2	AK	636	11.4	7.63	1	Alaska
3	3	AZ	400	18.2	8.31	1	Arizona
4	4	AR	480	18.7	9.41	1	Arkansas
5	5	CA	396	16.4	7.25	0	California
6	6	CO	309	12.1	6.75	0	Colorado
7	7	CT	237	10.8	8.04	0	Connecticut
8	8	DE	489	13.0	6.52	0	Delaware
9	9	DC	1244	18.4	8.41	0	District of Colu~
10	10	FL	540	16.6	8.29	1	Florida

```
# ... with 41 more rows
```

Numerical Summaries




```
crimestat %>% select(poverty, single, crime) %>% skim
```

```
Skim summary statistics
```

```
n obs: 51
```

```
n variables: 3
```

```
Variable type: numeric
```

variable	missing	complete	n	mean	sd	p0	p25	median	p75	p100	hist
crime	0	51	51	364.41	179.05	99.3	260.2	326.5	427.35	1244.4	
poverty	0	51	51	14.87	3.08	9.2	12.15	14.8	17.2	21.9	
single	0	51	51	7.69	1.61	4.48	6.75	7.63	8.5	11.59	

Modeling crime with poverty and single

Our main goal will be to build a linear regression model to predict **crime** using **poverty** and **single**.

We'll start by building an ordinary least squares model on the two predictors (after centering them, so that the intercept is meaningful) and looking at some diagnostics.

```
crimestat <- crimestat %>%  
  mutate(pov_c = poverty - mean(poverty),  
         single_c = single - mean(single))
```

Fitting an OLS model

Our first model mod1 using OLS

```
(mod1 <- lm(crime ~ pov_c + single_c, data = crimestat))
```

Call:

```
lm(formula = crime ~ pov_c + single_c, data = crimestat)
```

Coefficients:

(Intercept)	pov_c	single_c
364.41	16.11	23.84

```
confint(mod1)
```

	2.5 %	97.5 %
(Intercept)	318.296950	410.51481
pov_c	-3.218922	35.44816
single_c	-13.121152	60.80677

glance(mod1) and tidy(mod1)



```
      r.squared adj.r.squared   sigma statistic      p.value df
1  0.196879      0.1634156 163.771   5.883417 0.005184941   3

      logLik      AIC      BIC deviance df.residual
1 -330.8419 669.6837 677.411  1287405           48

      term      estimate std.error statistic      p.value
1 (Intercept) 364.40588 22.932525 15.890351 9.475916e-21
2      pov_c   16.11462  9.615642  1.675876 1.002655e-01
3    single_c   23.84281 18.384226  1.296917 2.008596e-01
```

Neither predictor meets our usual standard of having an estimate which is at least twice as large as the standard error.

Residual Plots for our model?



Potential Problems: States 9, 25 and maybe 2

Who are the outlier states?

Points 9, 25 and maybe 2 look like they could be problematic. Which states are those?

```
filter(crimestat, row_number() %in% c(9, 25, 2))
```

```
# A tibble: 3 x 9
```

	sid	state	crime	poverty	single	trump	state.full	pov_c
	<int>	<fct>	<dbl>	<dbl>	<dbl>	<int>	<fct>	<dbl>
1	2	AK	636	11.4	7.63	1	Alaska	-3.47
2	9	DC	1244	18.4	8.41	0	District of~	3.53
3	25	MS	278	21.9	11.4	1	Mississippi	7.03

```
# ... with 1 more variable: single_c <dbl>
```

Augmented Data Set with OLS results

```
crime_with_mod1 <- augment(mod1, crimestat)
head(crime_with_mod1, 3)
```

	sid	state	crime	poverty	single	trump	state.full	pov_c
1	1	AL	427.4	19.2	9.02	1	Alabama	4.327451
2	2	AK	635.8	11.4	7.63	1	Alaska	-3.472549
3	3	AZ	399.9	18.2	8.31	1	Arizona	3.327451


	single_c	.fitted	.se.fit	.resid	.hat
1	1.33117647	465.8801	39.84145	-38.48010	0.05918290
2	-0.05882353	307.0446	39.96271	328.75545	0.05954371
3	0.62117647	432.8371	34.99599	-32.93709	0.04566282

	.sigma	.cooks	.std.resid
1	165.4029	0.0012304477	-0.2422405
2	157.9444	0.0904293530	2.0699826
3	165.4310	0.0006759806	-0.2058720

Standardized Residuals vs. Fitted Values (code)

```
ggplot(crime_with_mod1, aes(x = .fitted, y = .std.resid,  
                             size = abs(.std.resid),  
                             col = -abs(.std.resid))) +  
  geom_label(aes(label = state)) +  
  guides(size = FALSE, col = FALSE) +  
  scale_color_continuous(low = "red", high = "black") +  
  theme_bw() +  
  labs(x = "Fitted Value from Linear Model",  
       y = "Standardized Residual")
```


Standardized Residuals vs. Fitted Values



Cook's Distance vs. |Standardized Residuals| (code)

```
ggplot(crime_with_mod1, aes(x = abs(.std.resid), y = .cooks,
                             size = .cooks,
                             col = .cooks < 0.2)) +
  geom_label(aes(label = state)) +
  guides(size = FALSE, col = FALSE) +
  scale_color_discrete() +
  theme_bw() +
  labs(x = "|Standardized Residual from Linear Model|",
       y = "Cook's Distance from Linear Model")
```

Cook's Distance vs. |Standardized Residuals|



What about just “robustifying” the standard errors of the coefficients?

Would Sandwich Estimation help for our original model?

from [David Freedman](#):

The “Huber Sandwich Estimator” (for which Peter Huber is not to be blamed) can be used to estimate the variance of the MLE (maximum likelihood estimate) when the underlying model is incorrect. If the model is nearly correct, so are the usual standard errors, and robustification is unlikely to help much. On the other hand, if the model is seriously in error, the sandwich may help on the variance side, but the parameters being estimated by the MLE are likely to be meaningless.

Sandwich estimation is mainly used to help address heteroscedasticity in linear regression, not so much with outliers. So, I doubt it will get us all the way to significance, but let's see. . .

Using `lmtest::coeftest` to get standard errors

```
mod1 <- lm(crime ~ pov_c + single_c, data = crimestat)

# requires lmtest package
coeftest(mod1)
```

t test of coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	364.4059	22.9325	15.8904	<2e-16 ***
pov_c	16.1146	9.6156	1.6759	0.1003
single_c	23.8428	18.3842	1.2969	0.2009

Signif. codes:

0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Using coeftest for Robust (Huber) Standard Errors

```
# requires lmtest and sandwich packages  
coeftest(mod1, vcov = sandwich)
```

t test of coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	364.406	22.248	16.3794	< 2e-16 ***
pov_c	16.115	10.898	1.4787	0.14574
single_c	23.843	13.948	1.7095	0.09382 .

Signif. codes:

0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Using coefci for Robust (Huber) Standard Errors

```
# requires lmtest and sandwich packages  
coefci(mod1, vcov = sandwich, level = 0.95)
```

	2.5 %	97.5 %
(Intercept)	319.673648	409.13812
pov_c	-5.796371	38.02561
single_c	-4.200619	51.88624

Bootstrapping Regression Coefficients

Bootstrapped Regression Coefficient Estimates

I'd be happier using bootstrapped estimates in this setting.

Source: statmethods.net link

```
# requires boot package
# build function to obtain regression weights
bs <- function(formula, data, indices) {
  d <- data[indices,] # allows boot to select sample
  fit <- lm(formula, data=d)
  return(coef(fit))
}

# now do R = 1000 replications
set.seed(432222)
results <- boot(data=crimestat, statistic=bs,
  R=1000, formula = crime ~ pov_c + single_c)
```

Bootstrapping Estimates with 1,000 replications

```
results
```

ORDINARY NONPARAMETRIC BOOTSTRAP

Call:

```
boot(data = crimestat, statistic = bs, R = 1000, formula = cri  
      pov_c + single_c)
```

Bootstrap Statistics :



	original	bias	std. error
t1*	364.40588	2.107998	22.63315
t2*	16.11462	1.235138	11.58621
t3*	23.84281	-1.224433	15.51142

Plots of Bootstrapped Estimates (next 3 slides)

```
plot(results, index = 1) # intercept  
plot(results, index = 2) # pov_c slope  
plot(results, index = 3) # single_c slope
```



Intercept Estimates (bootstrap)

Histogram of t





pov_c Slope Estimates (bootstrap)

Histogram of t



single_c Slope Estimates (bootstrap)

Histogram of t



Obtain 95% Confidence Intervals

```
boot.ci(results, type="bca", index=1) # intercept  
boot.ci(results, type="bca", index=2) # pov_c slope  
boot.ci(results, type="bca", index=3) # single_c slope
```

95% Bootstrap CI for Intercept

```
boot.ci(results, type="bca", index=1) # intercept
```

BOOTSTRAP CONFIDENCE INTERVAL CALCULATIONS

Based on 1000 bootstrap replicates

CALL :

```
boot.ci(boot.out = results, type = "bca", index = 1)
```

Intervals :

Level	BCa
-------	-----

95%	(328.9, 421.8)
-----	-----------------

Calculations and Intervals on Original Scale

95% Bootstrap CI for Slope of pov_c

```
boot.ci(results, type="bca", index=2) # pov_c slope
```

BOOTSTRAP CONFIDENCE INTERVAL CALCULATIONS

Based on 1000 bootstrap replicates

CALL :

```
boot.ci(boot.out = results, type = "bca", index = 2)
```

Intervals :

Level	BCa
-------	-----

95%	(-2.90, 41.12)
-----	-----------------

Calculations and Intervals on Original Scale

95% Bootstrap CI for Slope of `single_c`

```
boot.ci(results, type="bca", index=3) # single_c slope
```

BOOTSTRAP CONFIDENCE INTERVAL CALCULATIONS

Based on 1000 bootstrap replicates

CALL :

```
boot.ci(boot.out = results, type = "bca", index = 3)
```

Intervals :

Level	BCa
-------	-----

95%	(-11.19, 50.37)
-----	------------------

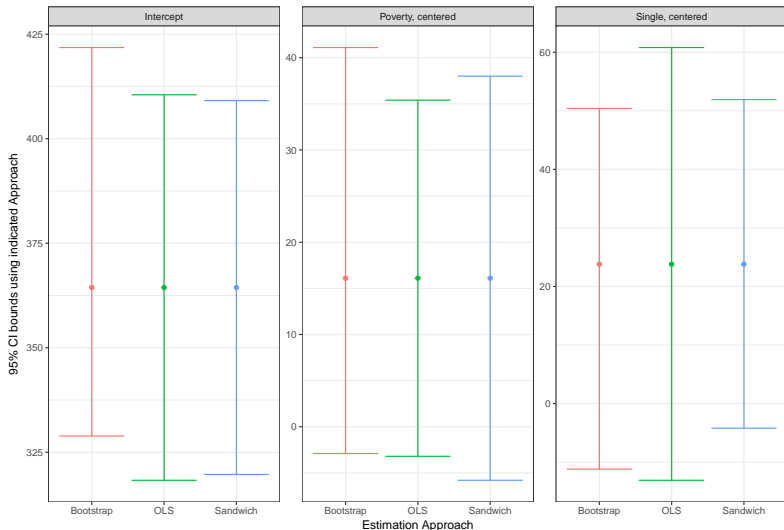
Calculations and Intervals on Original Scale

Standard, Sandwich and Bootstrapped 95% CIs for the Coefficients of our OLS model

Fitted OLS Model: $\text{crime} = 364.4 + 16.1 \times \text{pov_c} + 23.8 \times \text{single_c}$

	Fit	Intercept CI	pov_c CI	single_c CI
OLS		(318.3, 410.5)	(-3.2, 35.4)	(-13.1, 60.8)
OLS with sandwich		(319.7, 409.1)	(-5.8, 38.0)	(-4.2, 51.9)
OLS, bootstrapped		(328.9, 421.8)	(-2.9, 41.1)	(-11.2, 50.4)

Comparison Plot (code, next two slides)



Point Estimates from OLS: not varying here by Approach

Code for plot on prior slide (part 1)

```
res_class14 <- data_frame(  
  approach = c(rep("OLS",3), rep("Sandwich",3),  
               rep("Bootstrap",3)),  
  parameter = c(rep(c("Intercept", "Poverty, centered",  
                      "Single, centered"),3)),  
  estimate = c(rep(c(364.4, 16.1, 23.8),3)),  
  conf.low = c(318.3, -3.2, -13.1, 319.7, -5.8, -4.2,  
              328.9, -2.9, -11.2),  
  conf.high = c(410.5, 35.4, 60.8, 409.1, 38.0, 51.9,  
              421.8, 41.1, 50.4)  
)
```

Code for plot on prior slide (part 2)

```
ggplot(res_class14, aes(x = approach, y = estimate,  
                        col = approach)) +  
  geom_point() +  
  geom_errorbar(aes(ymin = conf.low, ymax = conf.high)) +  
  labs(x = "Estimation Approach",  
       y = "95% CI bounds using indicated Approach",  
       caption = "Point Estimates from OLS:  
not varying here by Approach") +  
  guides(col = FALSE) +  
  theme_bw() +  
  facet_wrap(~ parameter, scales = "free_y")
```

Good luck on the Quiz!

Due Monday at Noon.